

# Address Space Layout Randomization

---

Vitaly Shmatikov

# Reading Assignment

---

- ◆ Shacham et al. “On the effectiveness of address-space randomization” (CCS 2004).
- ◆ Optional:
  - PaX documentation (<http://pax.grsecurity.net/docs/>)
  - Bhatkar, Sekar, DuVarney. “Efficient techniques for comprehensive protection from memory error exploits” (Usenix Security 2005).

# Problem: Lack of Diversity

---

- ◆ Buffer overflow and **return-to-libc** exploits need to know the (virtual) address to hijack control
  - Address of attack code in the buffer
  - Address of a standard kernel library routine
- ◆ Same address is used on many machines
  - Slammer infected 75,000 MS-SQL servers using same code on every machine
- ◆ Idea: introduce **artificial diversity**
  - Make stack addresses, addresses of library routines, etc. unpredictable and different from machine to machine

# ASLR

---

- ◆ Address Space Layout Randomization
- ◆ Randomly choose base address of stack, heap, code segment
- ◆ Randomly pad stack frames and malloc() calls
- ◆ Randomize location of Global Offset Table
- ◆ Randomization can be done at compile- or link-time, or by rewriting existing binaries
  - Threat: attack repeatedly probes randomized binary

# PaX

---

- ◆ Linux kernel patch
- ◆ Goal: prevent execution of arbitrary code in an existing process's memory space
- ◆ Enable executable/non-executable memory pages
- ◆ Any section not marked as executable in ELF binary is non-executable by default
  - Stack, heap, anonymous memory regions
- ◆ Access control in `mmap()`, `mprotect()` prevents unsafe changes to protection state at runtime
- ◆ Randomize address space layout

# Non-Executable Pages in PaX

---

- ◆ In older x86, pages cannot be directly marked as non-executable
- ◆ PaX marks each page as “non-present” or “supervisor level access”
  - This raises a page fault on every access
- ◆ Page fault handler determines if the fault occurred on a data access or instruction fetch
  - Instruction fetch: log and terminate process
  - Data access: unprotect temporarily and continue

# mprotect() in PaX

---

- ◆ mprotect() is a Linux kernel routine for specifying desired protections for memory pages
- ◆ PaX modifies mprotect() to prevent:
  - Creation of executable anonymous memory mappings
  - Creation of executable and writable file mappings
  - Making executable, read-only file mapping writable
    - Except when relocating the binary
  - Conversion of non-executable mapping to executable

# Access Control in PaX

## mprotect()

- ◆ In standard Linux kernel, each memory mapping is associated with permission bits
  - `VM_WRITE`, `VM_EXEC`, `VM_MAYWRITE`, `VM_MAYEXEC`
    - Stored in the `vm_flags` field of the `vma` kernel data structure
    - 16 possible write/execute states for each memory page
- ◆ PaX makes sure that the same page cannot be writable AND executable at the same time
  - Ensures that the page is in one of the 4 “good” states
    - `VM_MAYWRITE`, `VM_MAYEXEC`, `VM_WRITE | VM_MAYWRITE`, `VM_EXEC | VM_MAYEXEC`
  - Also need to ensure that attacker cannot make a region executable when mapping it using `mmap()`



# PaX ASLR

---

- ◆ User address space consists of three areas
  - Executable, mapped, stack
- ◆ Base of each area shifted by a random “delta”
  - Executable: 16-bit random shift (on x86)
    - Program code, uninitialized data, initialized data
  - Mapped: 16-bit random shift
    - Heap, dynamic libraries, thread stacks, shared memory
    - Why are only 16 bits of randomness used?
  - Stack: 24-bit random shift
    - Main user stack

# PaX RANDUSTACK

---

- ◆ Responsible for randomizing userspace stack
- ◆ Userspace stack is created by the kernel upon each `execve()` system call
  - Allocates appropriate number of pages
  - Maps pages to process's virtual address space
    - Userspace stack is usually mapped at `0xBFFFFFFF`, but PaX chooses a random base address
- ◆ In addition to base address, PaX randomizes the range of allocated memory

# PaX RANDKSTACK

---

- ◆ Linux assigns two pages of kernel memory for each process to be used during the execution of system calls, interrupts, and exceptions
- ◆ PaX randomizes each process's kernel stack pointer before returning from kernel to userspace
  - 5 bits of randomness
- ◆ Each system call is randomized differently
  - By contrast, user stack is randomized once when the user process is invoked for the first time

# PaX RANDMMAP

---

- ◆ Linux heap allocation: `do_mmap()` starts at the base of the process's unmapped memory and looks for the first unallocated chunk which is large enough
- ◆ PaX: add a random `delta_mmap` to the base address before looking for new memory
  - 16 bits of randomness

# PaX RANDEXEC

---

- ◆ Randomizes location of ELF binaries in memory
- ◆ Problem if the binary was created by a linker which assumed that it will be loaded at a fixed address and omitted relocation information
  - PaX maps the binary to its normal location, but makes it non-executable + creates an executable mirror copy at a random location
  - Access to the normal location produces a page fault
  - Page handler redirects to the mirror “if safe”
    - Looks for “signatures” of return-to-libc attacks and may result in false positives

# Base-Address Randomization

---

- ◆ Only the base address is randomized
  - **Layouts** of stack and library table remain the same
  - Relative distances between memory objects are not changed by base address randomization
- ◆ To attack, it's enough to guess the base shift
- ◆ A 16-bit value can be guessed by brute force
  - Try  $2^{15}$  (on average) overflows with different values for addr of known library function – how long does it take?
    - Shacham et al. attacked Apache with return-to-libc
    - usleep() is used (why?)
  - If address is wrong, target will simply crash

# ASLR in Windows

---

- ◆ Vista and Server 2008
- ◆ Stack randomization
  - Find  $N^{\text{th}}$  hole of suitable size (N is a 5-bit random value), then random word-aligned offset (9 bits of randomness)
- ◆ Heap randomization: 5 bits
  - Linear search for base + random 64K-aligned offset
- ◆ EXE randomization: 8 bits
  - Preferred base + random 64K-aligned offset
- ◆ DLL randomization: 8 bits
  - Random offset in DLL area; random loading order

# Bypassing Windows ASLR

---

- ◆ Implementation uses randomness improperly, thus distribution of heap bases is biased
  - Ollie Whitehouse's paper (Black Hat 2007)
  - Makes guessing a valid heap address easier
- ◆ When attacking browsers, may be able to insert arbitrary objects into the victim's heap
  - Executable JavaScript code, plugins, Flash, Java applets, ActiveX and .NET controls...
- ◆ **Heap spraying**
  - Stuff heap with large objects and multiple copies of attack code (how does this work?)



# Example: Java Heap Spraying

[See Sotirov & Dowd]

- ◆ JVM makes all of its allocated memory RWX: readable, writeable, executable (why?)
  - Yay! DEP now goes out the window...
- ◆ 100MB applet heap, randomized base in a predictable range
  - 0x20000000 through 0x25000000
- ◆ Use a Java applet to fill the heap with (almost) 100MB of NOP sleds + attack code
- ◆ Use your favorite memory exploit to transfer control to 0x25A00000 (why does this work?)

# Information Leaks Break ASLR

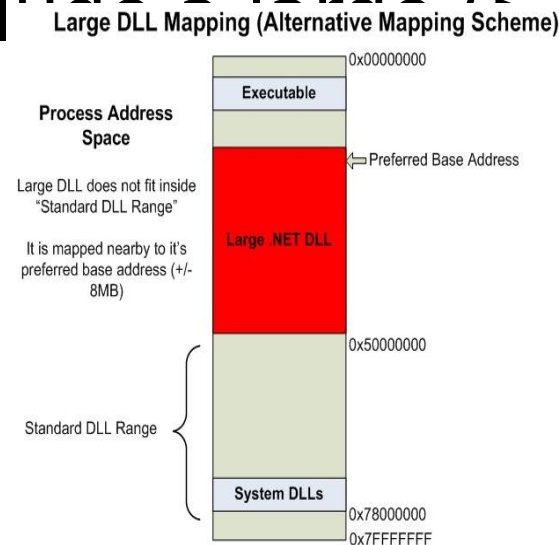
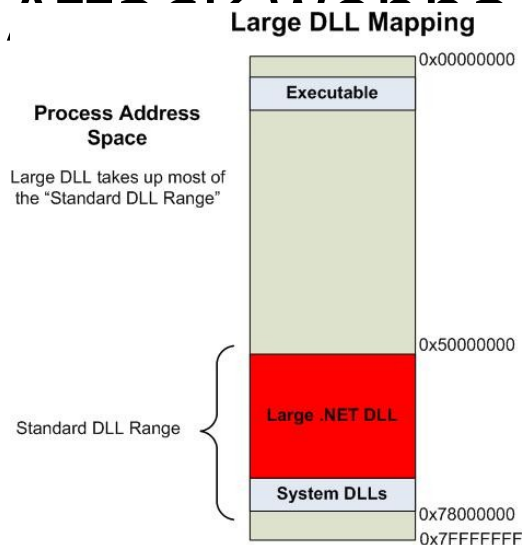
[See Sotirov & Dowd]

- ◆ User-controlled .NET objects are not RWX
- ◆ But JIT compiler generates code in RWX memory
  - Can overwrite this code or “return” to it out of context
  - But ASLR hides location of generated code stubs...
  - Call `MethodHandle.GetFunctionPointer()` ... .NET itself will tell you where the generated code lives!
- ◆ ASLR is often defeated by information leaks
  - Pointer betrays an object’s location in memory
    - For example, a pointer to a static variable reveals DLL’s location... for all processes on the system! (why?)
  - Pointer to a frame object betrays the entire stack

# .NET Address Space Spraying

[See Sotirov & Dowd]

- ◆ Webpage may embed .NET DLLs
  - No native code, only IL bytecode
  - Run in sandbox, thus no user warning (unlike ActiveX)
  - Mandatory base randomization when loaded
- ◆ Attack which page includes large (>100MB)



# Dealing with Large Attack

## DLLs

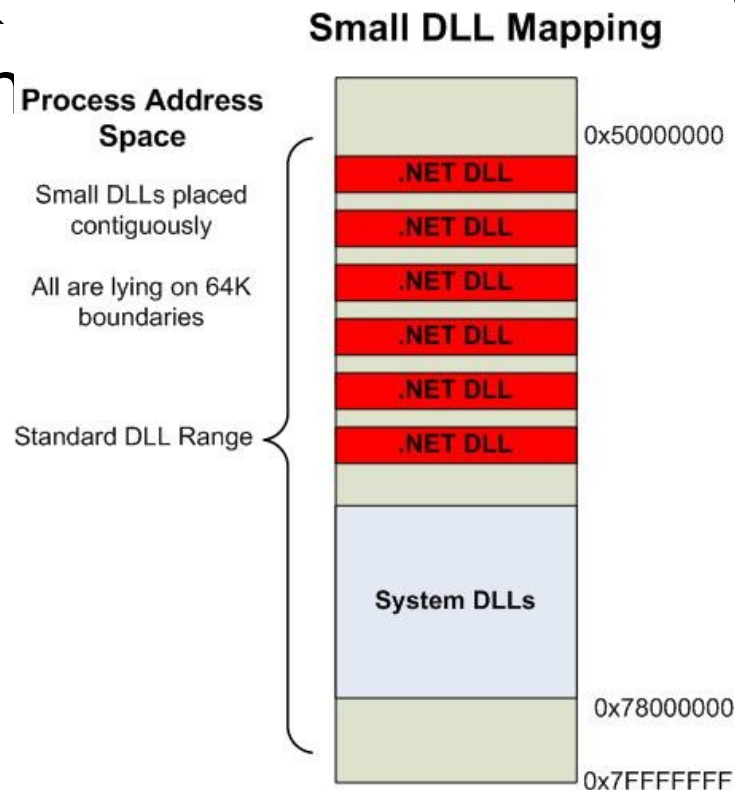
[See Sotirov & Dowd]

- ◆ 100MB is a lot for the victim to download!
- ◆ Solution 1: binary padding
  - Specify a section with a very large VirtualSize and very small SizeOfRawData – will be 0-padded when mapped
  - On x86, equivalent to `add byte ptr [eax], al` - NOP sled!
    - Only works if EAX points to a valid, writeable address
- ◆ Solution 2: compression
  - gzip content encoding
    - Great compression ratio, since content is mostly NOPs
  - Browser will unzip on the fly

# Spraying with Small DLLs

[See Sotirov & Dowd]

- ◆ Attack webpage includes many small DLL binaries
- ◆ Large chunk of address space will be sprayed with



# Turning Off ASLR Entirely

[See Sotirov & Dowd]

- ◆ Any DLL may “opt out” of ASLR
  - Choose your own ImageBase, unset `IMAGE_DLL_CHARACTERISTICS_DYNAMIC_BASE` flag
- ◆ Unfortunately, ASLR is enforced on IL-only DLL
- ◆ How does the loader know a binary is IL-only?

```
if( ( (pCORHeader->MajorRuntimeVersion > 2) ||  
(pCORHeader->MajorRuntimeVersion == 2 && pCORHeader->MinorRuntimeVersion  
>= 5) ) &&  
(pCORHeader->Flags & COMIMAGE_FLAGS_ILONLY) )  
{  
    pImageControlArea->pBinaryInfo->pHeaderInfo->bFlags |= PINFO_IL_ONLY_IMAGE;  
    ...  
}
```

Set version in the header to anything below 2.5  
ASLR will be disabled for this binary!

# Bypassing IL Protections

[Dowd & Sotirov, PacSec 2008]

- ◆ Embedded .NET DLLs are expected to contain IL bytecode only - many protection features
  - Verified prior to JIT compilation and at runtime, DEP
  - Makes it difficult to write effective shellcode
- ◆ ... enabled by a single global variable
  - `mscorwks!s_eSecurityState` must be set to 0 or 2
  - Does `mscorwks` participate in ASLR? **No!**
- ◆ Similar: disable Java bytecode verification
  - JVM does not participate in ASLR, either
  - To disable runtime verification, traverse the stack and set NULL protection domain for current method

# Ideas for Better Randomization (1)

---

## ◆ 64-bit addresses

- At least 40 bits available for randomization
  - Memory pages are usually between 4K and 4M in size
- Brute-force attack on 40 bits is not feasible

## ◆ Does more frequent randomization help?

- ASLR randomizes when a process is created
- Alternative: re-randomize address space while brute-force attack is still in progress
  - E.g., re-randomize non-forking process after each crash (recall that unsuccessful guesses result in target's crashing)
- This does not help much (why?)



# Ideas for Better Randomization (2)

---

- ◆ Randomly re-order entry points of library functions
  - Finding address of one function is no longer enough to compute addresses of other functions
    - What if attacker finds address of `system()`?
- ◆ ... at compile-time
  - Access to source, thus no virtual memory constraints; can use more randomness (any disadvantages?)
- ◆ ... or at run-time
  - How are library functions shared among processes?
  - How does normal code find library functions?

# Comprehensive Randomization (1)

[Bhatkar et al.]

## ◆ Function calls

- Convert all functions to function pointers and store them in an array
- Reorder functions within the binary
- Allocation order of arguments is randomized for each function call

## ◆ Indirect access to all static variables

- Accessed only via pointers stored in read-only memory
- Addresses chosen randomly at execution start

# Comprehensive Randomization (2)

[Bhatkar et al.]

- ◆ Locations of stack-allocated objects randomized continuously during execution
  - Separate shadow stack for arrays
  - Each array surrounded by inaccessible memory regions
- ◆ Insert random stack gap when a function is called
  - Can be done right before a function is called, or at the beginning of the called function (what's the difference?)
- ◆ Randomize heap-allocated objects
  - Intercepts malloc() calls and requests random amount of additional space

# Comprehensive Randomization (3)

[Bhatkar et al.]

- ◆ Randomize base of stack at program start
- ◆ Shared DLLs (see any immediate issues?)
- ◆ Procedure Linkage Table/Global Offset Table
- ◆ setjmp/longjmp require special handling
  - Must keep track of context (e.g., shadow stack location)

# Summary

---

- ◆ Randomness is a potential defense mechanism
  - ◆ Many issues for proper implementation
  - ◆ Serious limitations on 32-bit architecture
    - "Thus, on 32-bit systems, runtime randomization cannot provide more than 16-20 bits of entropy"
- Shacham et al.