



# Systems and Internet Infrastructure Security

Network and Security Research Center  
Department of Computer Science and Engineering  
Pennsylvania State University, University Park PA

## ***Return-oriented Programming***

*Trent Jaeger*

*Systems and Internet Infrastructure Security (SIIS) Lab  
Computer Science and Engineering Department  
Pennsylvania State University*

# Buffer Overflow

EIP

```
                                _start:
0x0804321      call main

                                int main() {
                                char buf[8];
0x0804480      gets(buf);
                                printf("You typed: %s", buf);
0x0804484      }
```


# Buffer Overflow

EIP

```
0x0804321  _start:
             call main

             int main() {
                 char buf[8];
0x0804480     gets(buf);
                 printf("You typed: %s", buf);
0x0804484     }
```

ESP




# Buffer Overflow

EIP

```
0x0804321  _start:
             call main

             int main() {
                 char buf[8];
0x0804480  gets(buf); 0x12345678 0x90abcdef 0xbf000004
             printf("You typed: %s", buf);
0x0804484  }
```

ESP




# Buffer Overflow

EIP

```
0x0804321  _start:
             call main

             int main() {
                 char buf[8];
0x0804480  gets(buf); 0x12345678 0x90abcdef 0xbf000004
             printf("You typed: %s", buf);
0x0804484  }
```

ESP



...


# Buffer Overflow

EIP

```
0x0804321  _start:
             call main

             int main() {
                 char buf[8];
0x0804480  gets(buf); 0x12345678 0x90abcdef 0xbf000004
             printf("You typed: %s", buf);
0x0804484  }
```

ESP




# Buffer Overflow

EIP

```
0x0804321  _start:
             call main

             int main() {
                 char buf[8];
0x0804480  gets(buf); 0x12345678 0x90abcdef 0xbf000004
             printf("You typed: %s", buf);
0x0804484  }
```

ESP



...


# Buffer Overflow

EIP

```
      _start:
0x0804321      call main

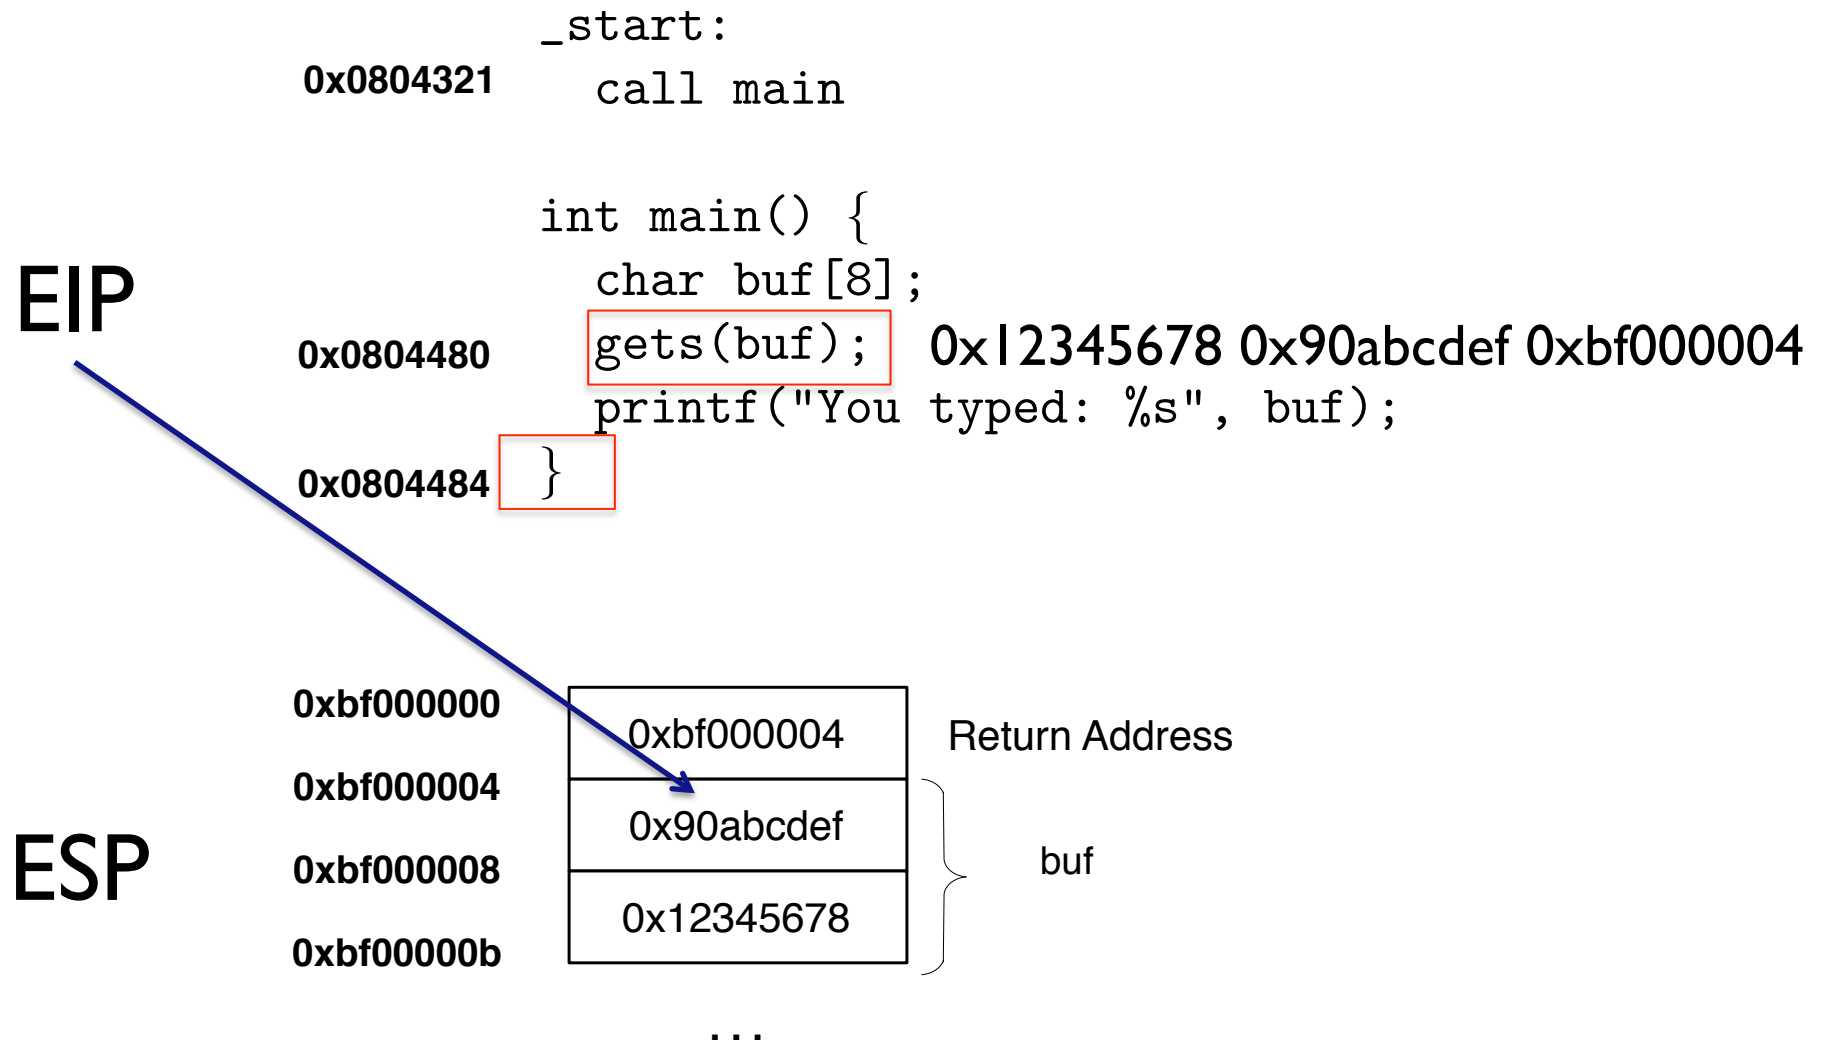
      int main() {
          char buf[8];
0x0804480      gets(buf);    0x12345678 0x90abcdef 0xbf000004
          printf("You typed: %s", buf);
0x0804484      }
```

ESP

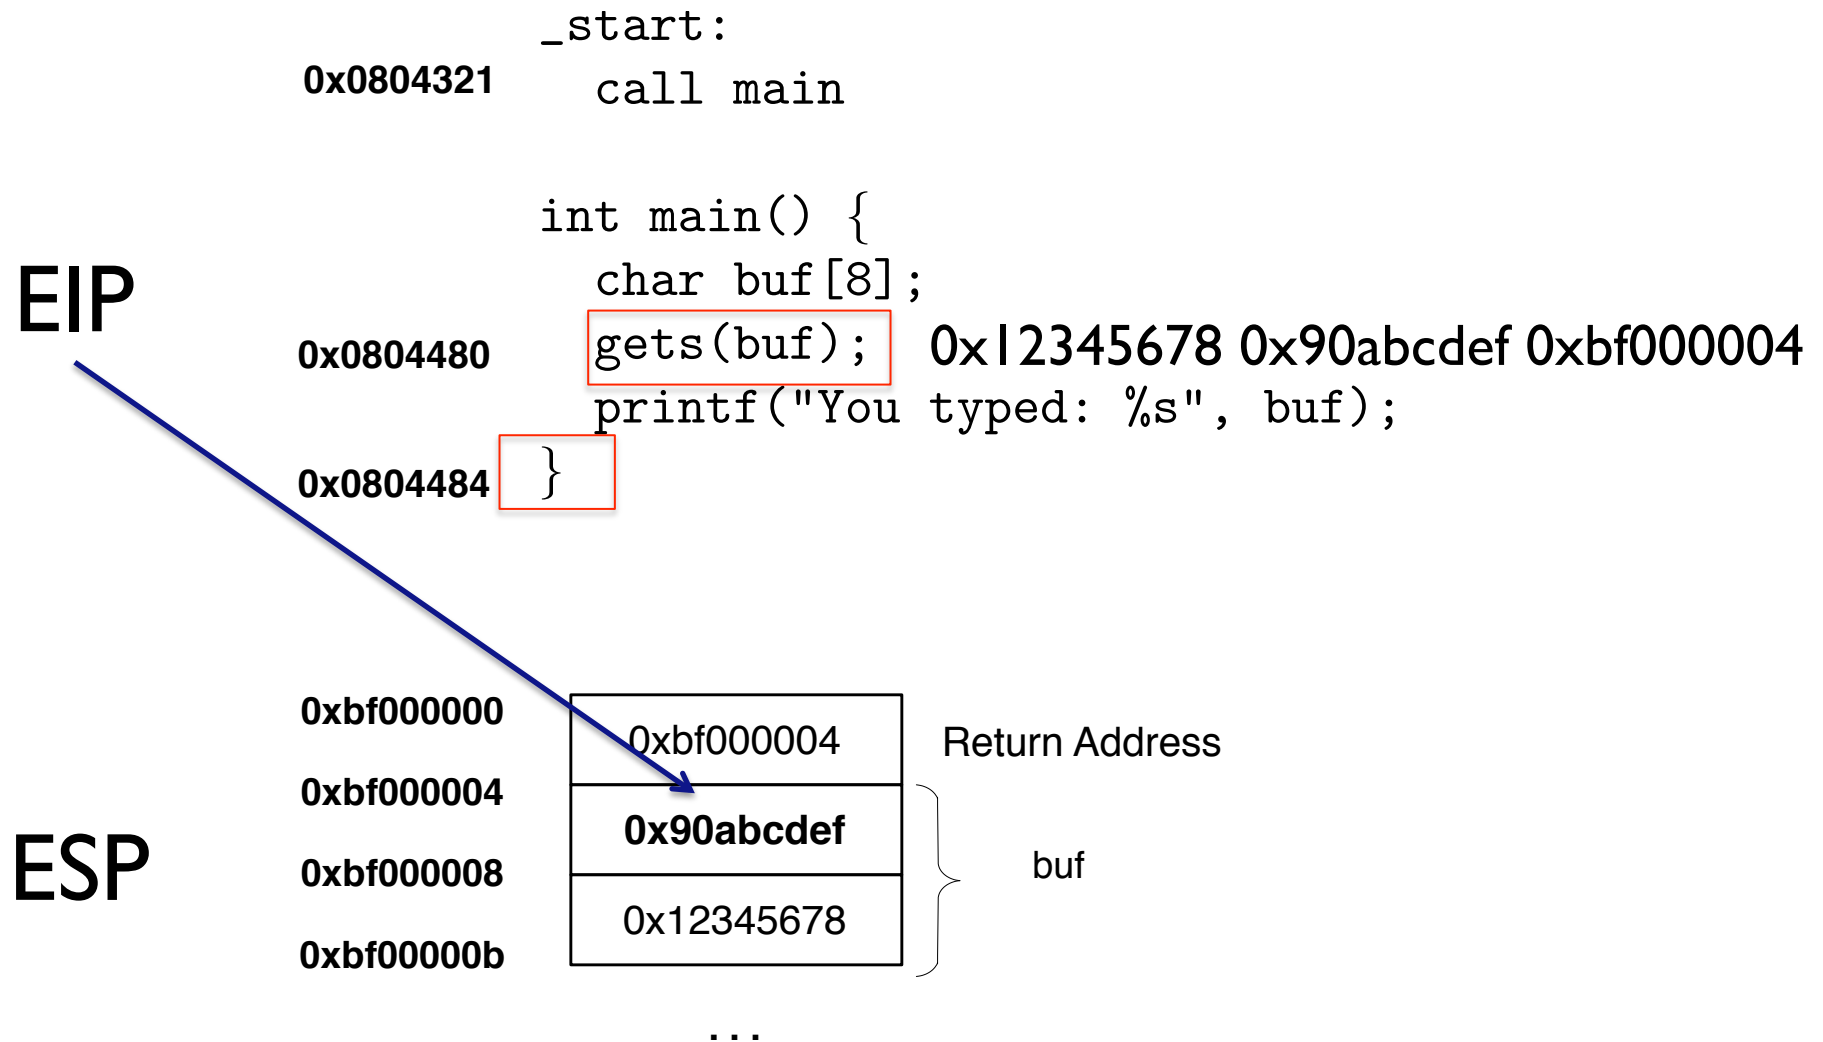




# Buffer Overflow




# Buffer Overflow




# Buffer Overflow Defense

- W xor X
  - Pages marked write can't be executed
- Return-to-libc




- Randomize bases of memory regions
  - ▶ Stack (Thwarts traditional stack overflow)
  - ▶ Mmap (Thwarts return-to-libc)
  - ▶ Brk (Heap – Thwarts traditional heap overflow)
  - ▶ Exec (Program binary)
- Not enabled by default



# Anatomy of Control Flow Attacks

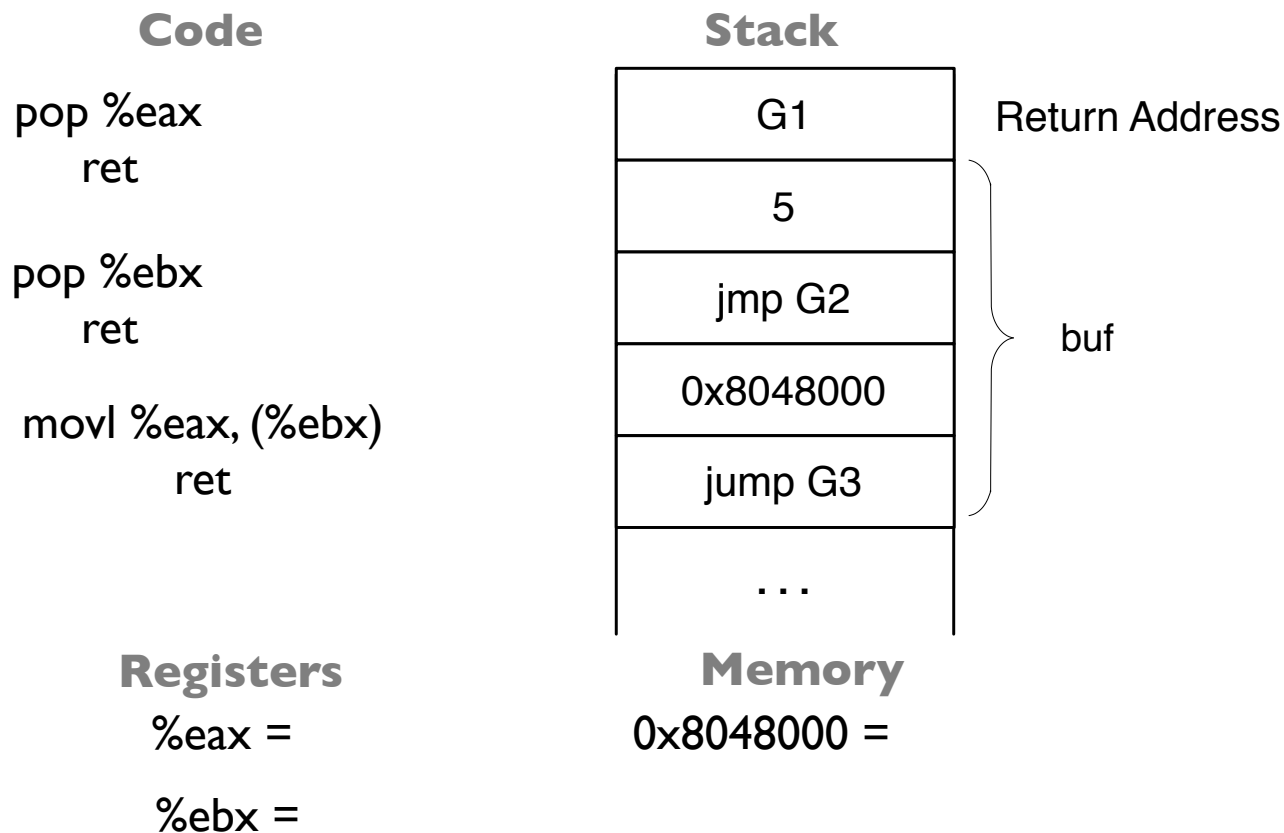
- Two steps
- First, the attacker changes the control flow of the program
  - ▶ In buffer overflow, overwrite the return address on the stack
  - ▶ What are the ways that this can be done?
- Second, the attacker uses this change to run code of their choice
  - ▶ In buffer overflow, inject code on stack
  - ▶ What are the ways that this can be done?



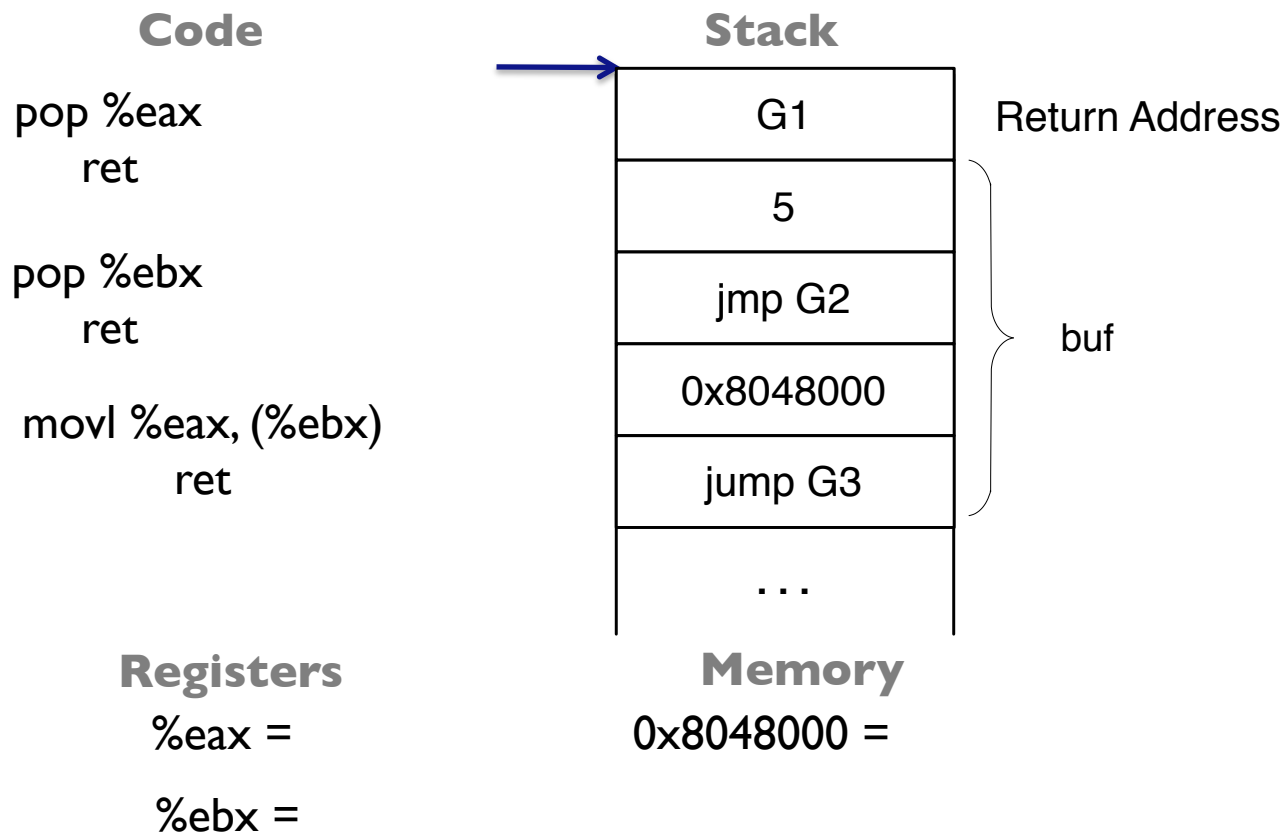
# Return-oriented Programming

- General approach to control flow attacks
- Demonstrates how general the two steps of a control flow attack can be
- First, change program control flow
  - In any way
- Then, run any code of attackers' choosing, including the code in the existing program

- Use ESP as program counter
  - E.g., Store 5 at address 0x8048000
    - without introducing new code

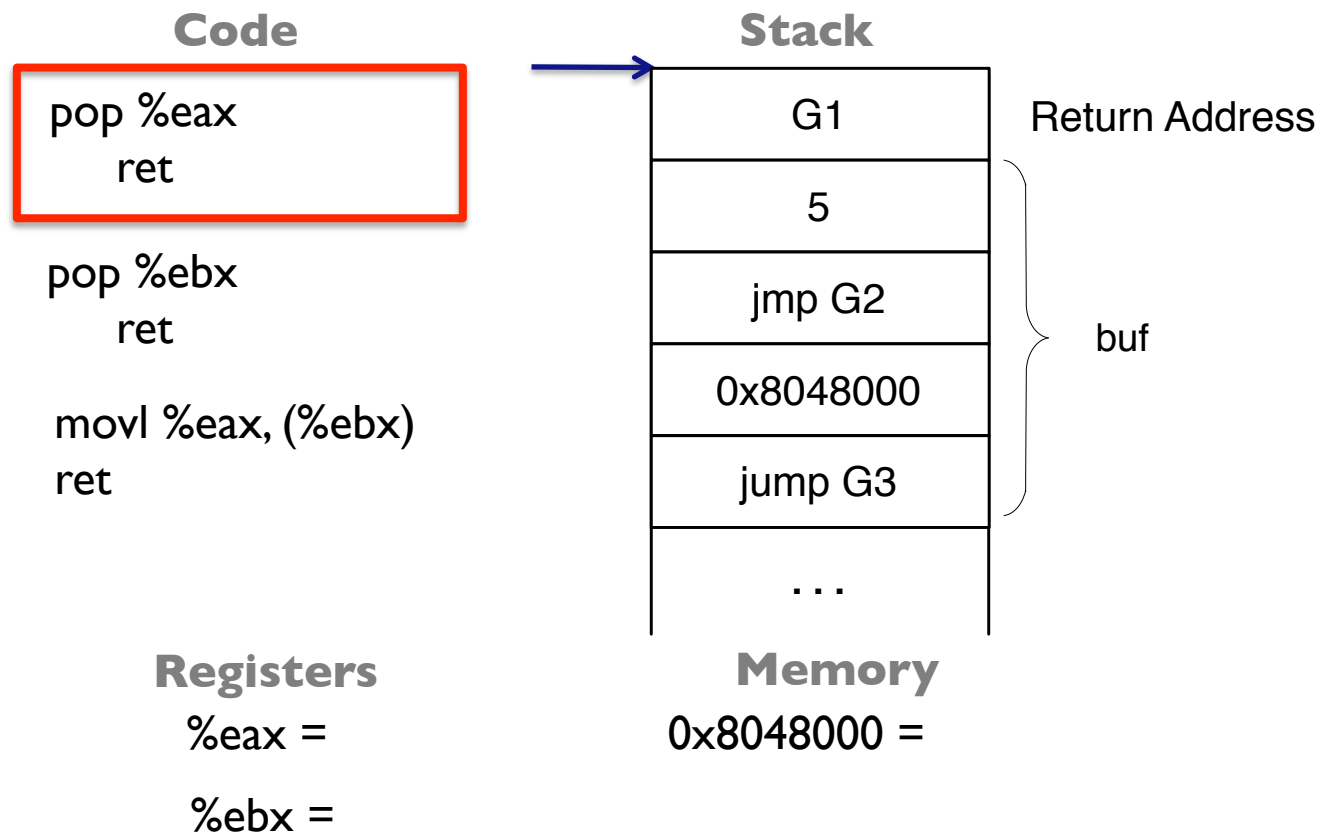


- Use ESP as program counter
  - E.g., Store 5 at address 0x8048000
    - without introducing new code

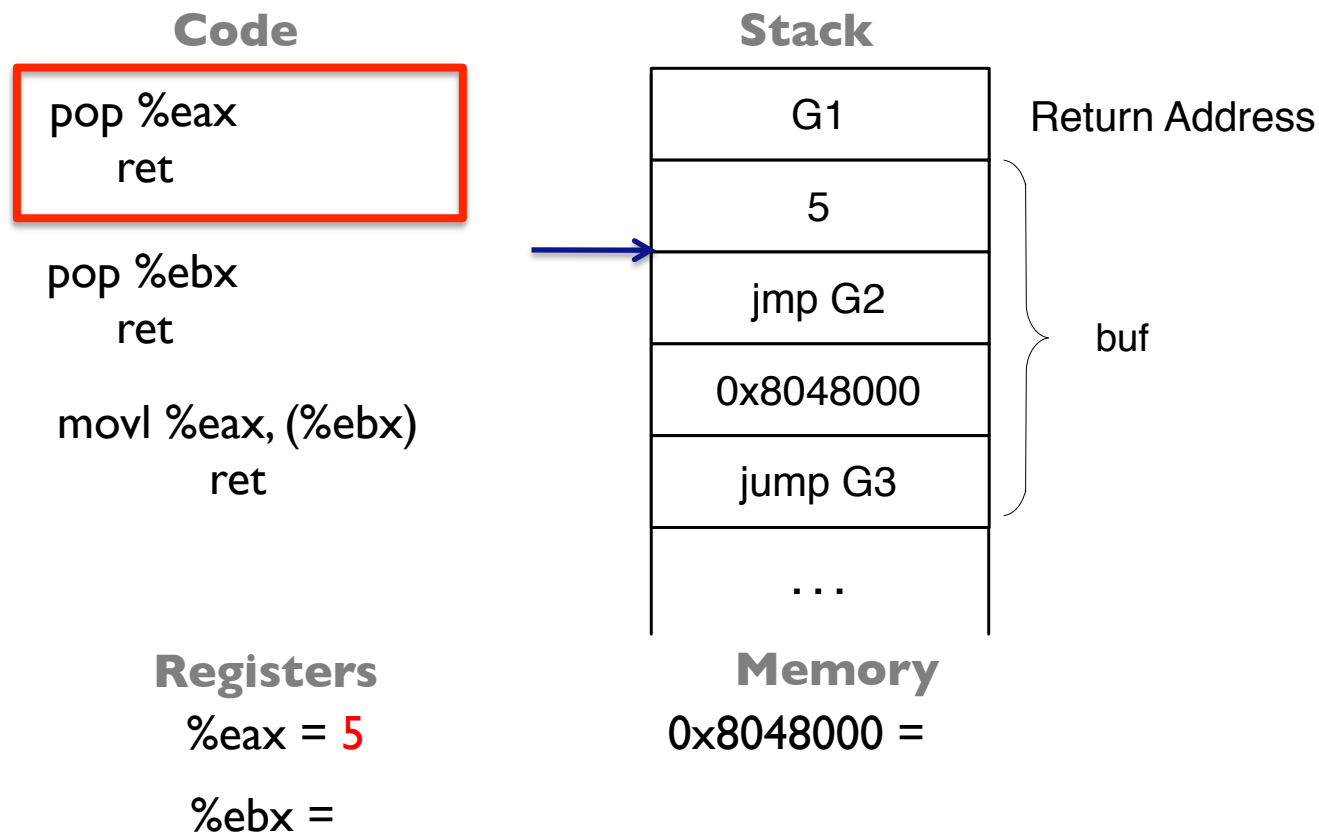




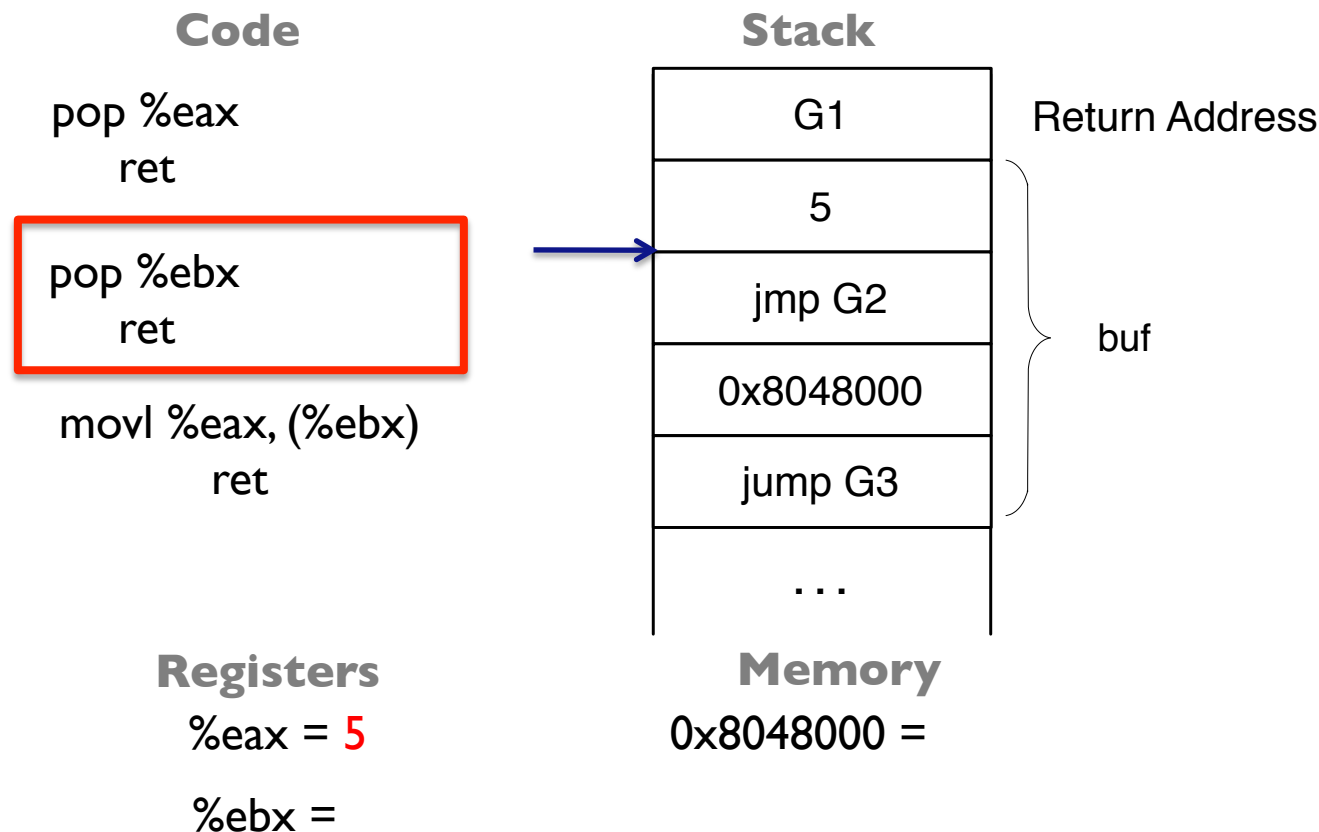
- Use ESP as program counter
  - E.g., Store 5 at address 0x8048000
    - without introducing new code



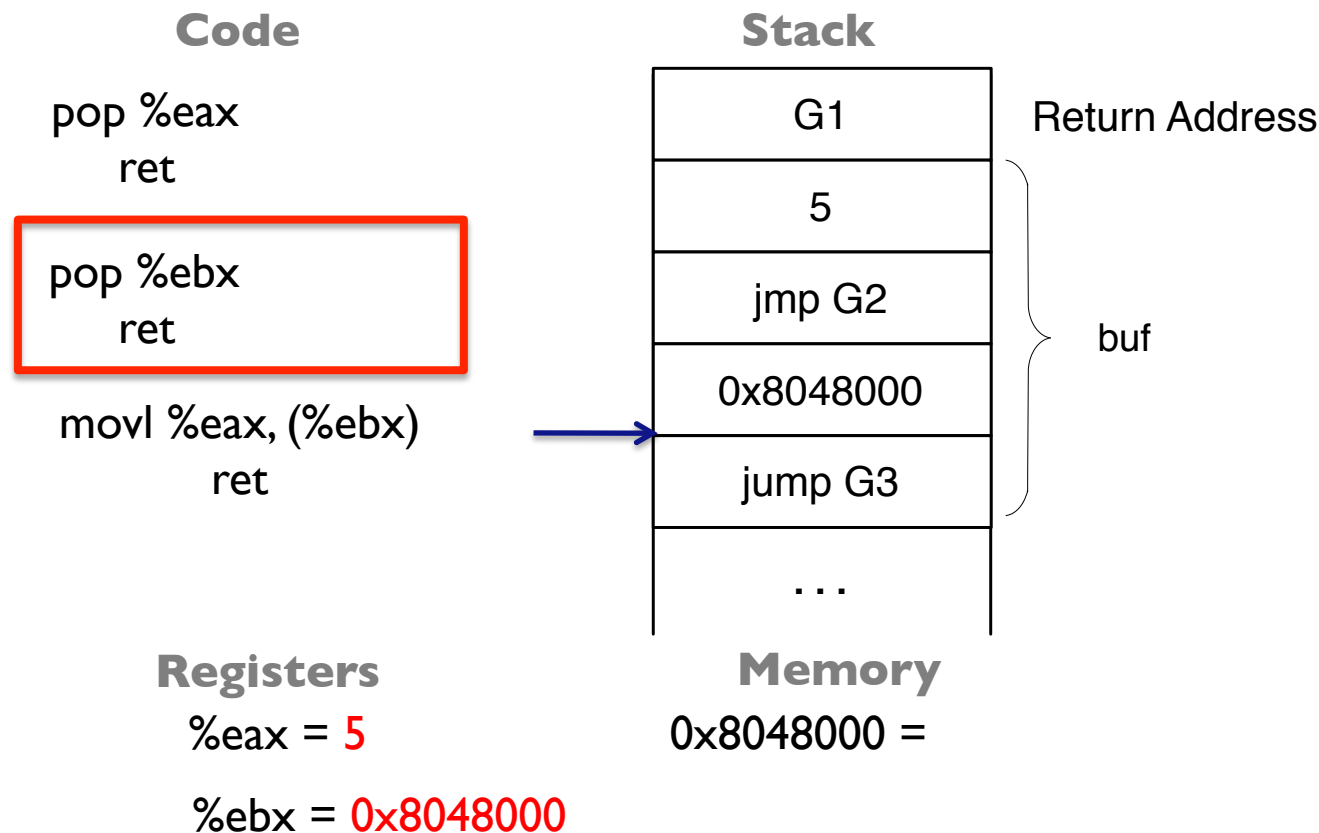
- Use ESP as program counter
  - E.g., Store 5 at address 0x8048000
    - without introducing new code



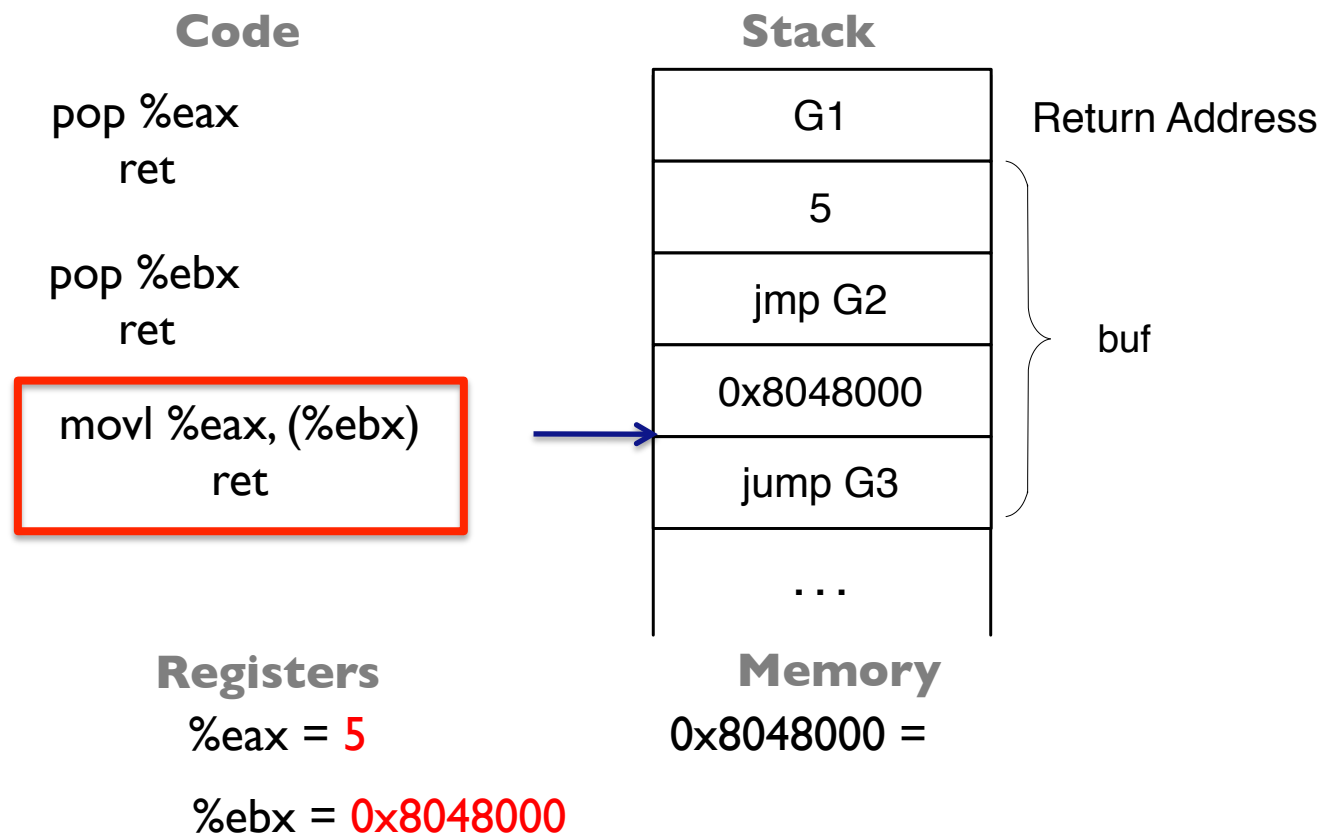
- Use ESP as program counter
  - E.g., Store 5 at address 0x8048000
    - without introducing new code



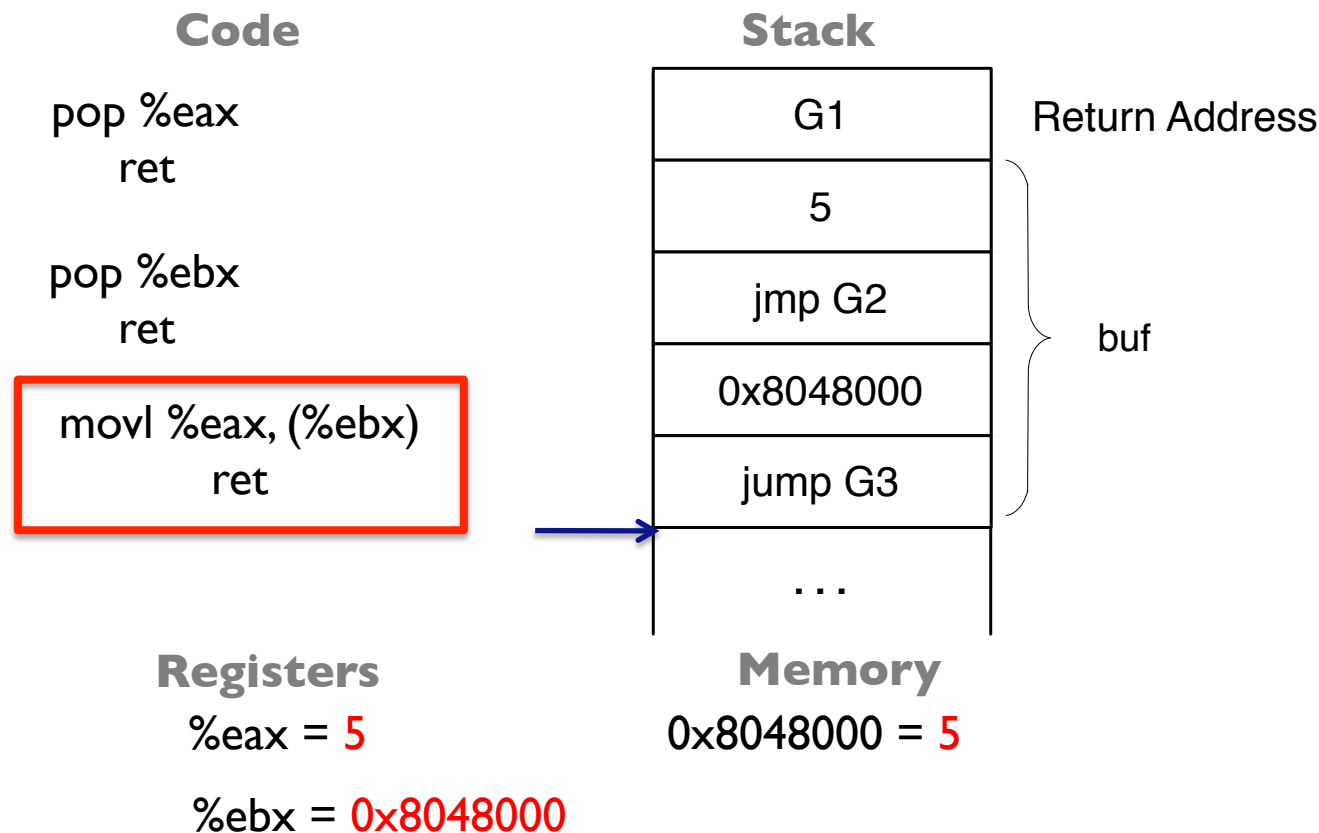
- Use ESP as program counter
  - E.g., Store 5 at address 0x8048000
    - without introducing new code



- Use ESP as program counter
  - E.g., Store 5 at address 0x8048000
    - without introducing new code



- Use ESP as program counter
  - E.g., Store 5 at address 0x8048000
    - without introducing new code




- How can an adversary make this happen?

## Return-oriented Programming: Exploitation without Code Injection

Erik Buchanan, Ryan Roemer, Stefan Savage, Hovav Shacham  
University of California, San Diego

# Return-oriented Programming


Bad code versus bad behavior



Problem: this implication is false!



any sufficiently large program codebase




arbitrary attacker computation and behavior,  
*without* code injection

(in the absence of control-flow integrity)

- ▶ Divert control flow of exploited program into libc code
  - ▶ `system()`, `printf()`,
- ▶ No code injection required
- ▶ Perception of return-into-libc: limited, easy to defeat
  - ▶ Attacker cannot execute arbitrary code
  - ▶ Attacker relies on contents of libc — remove `system()`?
- ▶ We show: this perception is *false*.

# ROP vs. Return-to-libc

attacker control of stack




arbitrary attacker computation and behavior  
via return-into-libc techniques

(given any sufficiently large codebase to draw on)


- ▶ Need control of memory around %esp
- ▶ Rewrite stack:
  - ▶ Buffer overflow on stack
  - ▶ Format string vuln to rewrite stack contents
- ▶ Move stack:
  - ▶ Overwrite saved frame pointer on stack;  
on leave/ret, move %esp to area under attacker control
  - ▶ Overflow function pointer to a register spring for %esp:
    - ▶ set or modify %esp from an attacker-controlled register
    - ▶ then return

# Machine Instructions



- ▶ Instruction pointer (%eip) determines which instruction to fetch & execute
- ▶ Once processor has executed the instruction, it automatically increments %eip to next instruction
- ▶ Control flow by changing value of %eip

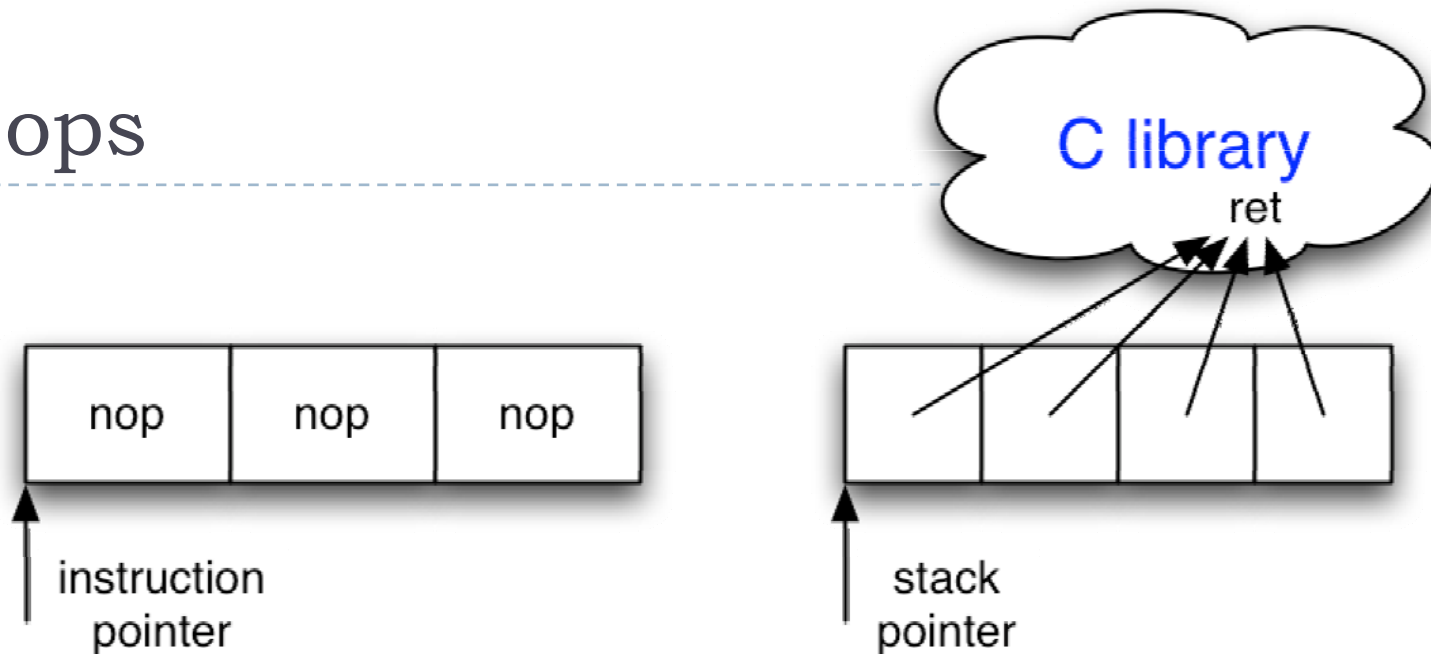
# ROP Execution



- ▶ *Stack pointer* (%esp) determines which instruction sequence to fetch & execute
- ▶ Processor doesn't automatically increment %esp; — but the "ret" at end of each instruction sequence does

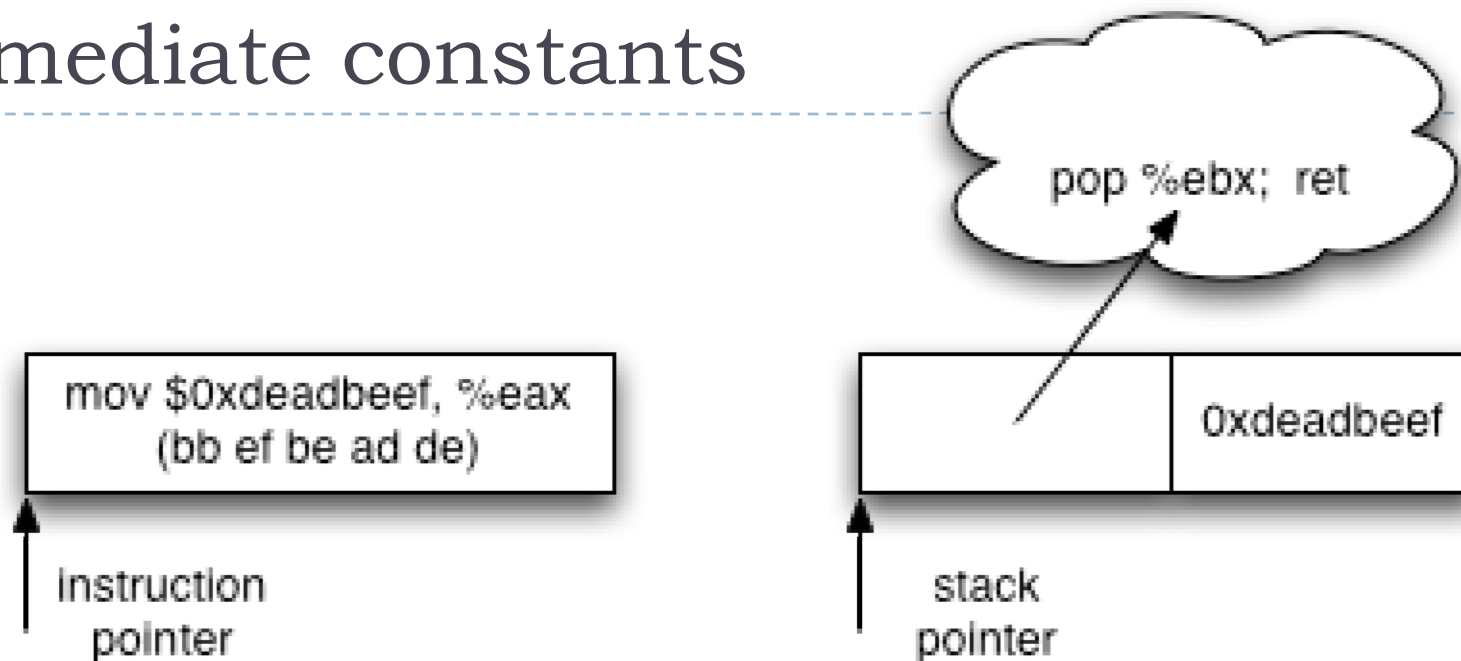
# Building ROP Functionality

## No-ops



- ▶ No-op instruction does nothing but advance %eip
- ▶ Return-oriented equivalent:
  - ▶ point to return instruction
  - ▶ advances %esp
- ▶ Useful in nop sled

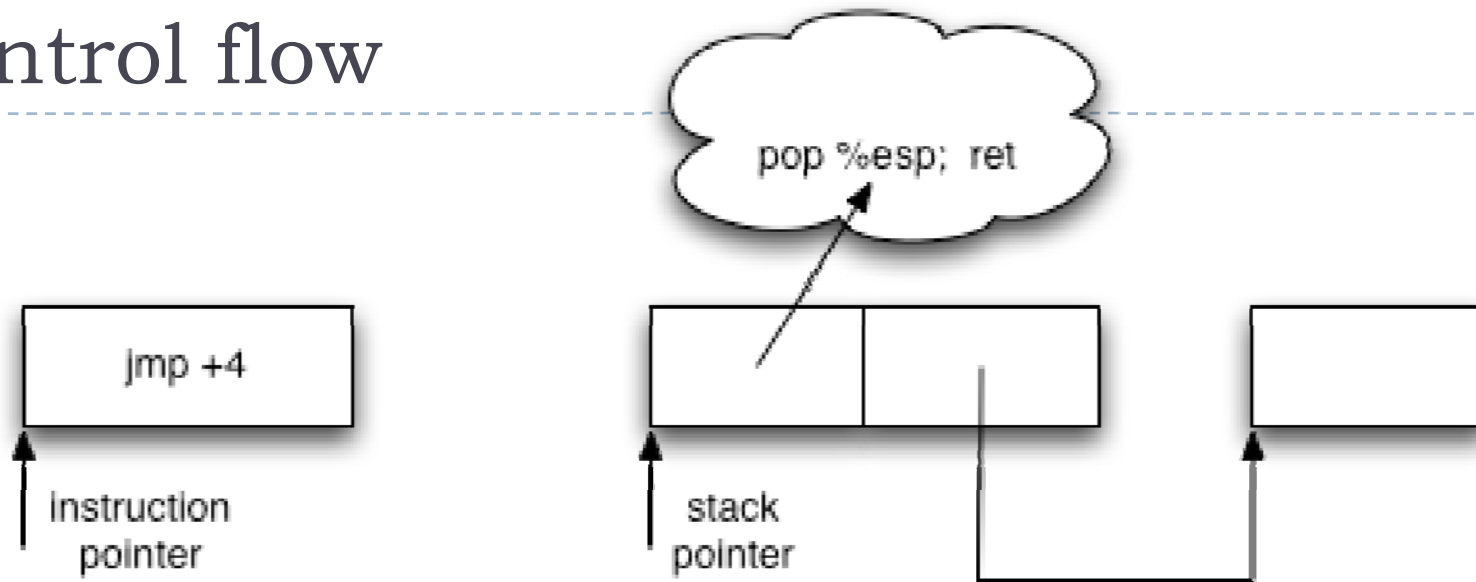
## Immediate constants



- ▶ Instructions can encode constants
- ▶ Return-oriented equivalent:
  - ▶ Store on the stack;
  - ▶ Pop into register to use

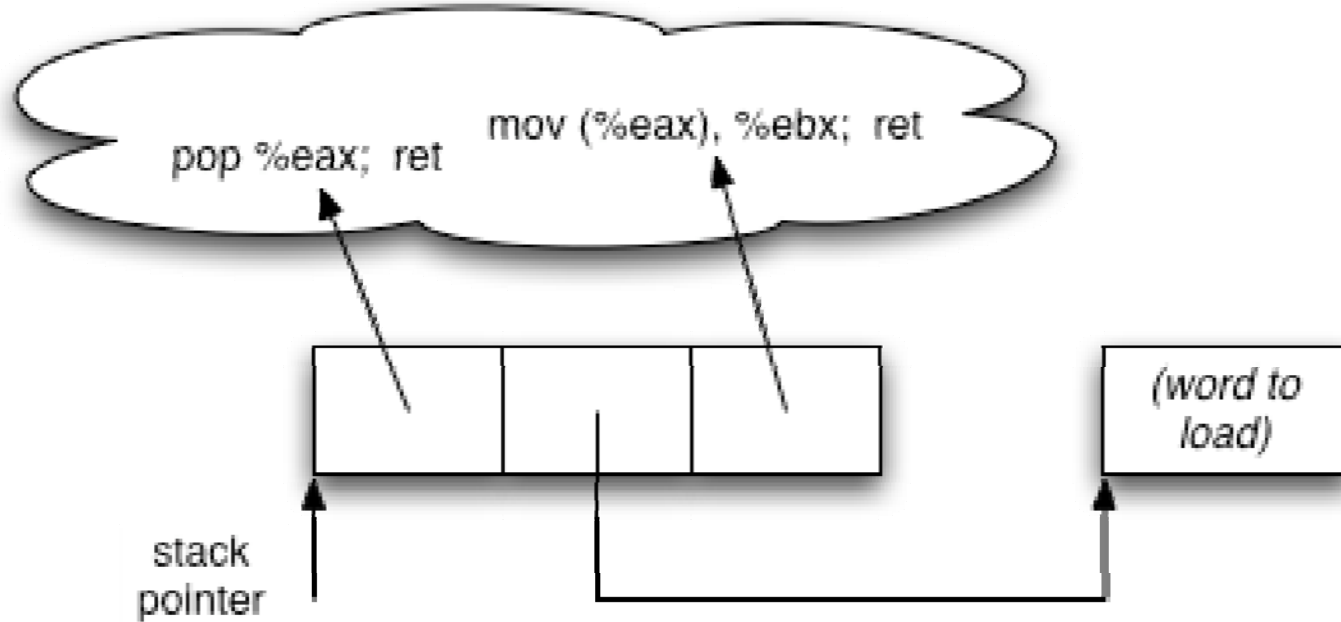


## Control flow



- ▶ Ordinary programming:
  - ▶ (Conditionally) set %eip to new value
- ▶ Return-oriented equivalent:
  - ▶ (Conditionally) set %esp to new value

## Gadgets: multiple instruction sequences



- ▶ Sometimes more than one instruction sequence needed to encode logical unit
- ▶ Example: load from memory into register:
  - ▶ Load address of source word into `%eax`
  - ▶ Load memory at `(%eax)` into `%ebx`

## Finding instruction sequences

---

- ▶ Any instruction sequence ending in “ret” is useful — could be part of a gadget
- ▶ **Algorithmic problem:** recover all sequences of valid instructions from libc that end in a “ret” insn
- ▶ Idea: at each ret (c3 byte) look back:
  - ▶ are preceding  $i$  bytes a valid length- $i$  insn?
  - ▶ recurse from found instructions
- ▶ Collect instruction sequences in a trie

## Return-oriented programming on SPARC

---

- ▶ Use Solaris 10 libc: 1.3 MB
- ▶ New techniques:
  - ▶ Use instruction sequences that are *suffixes* of real functions
  - ▶ Dataflow within a gadget:
    - ▶ Use structured dataflow to dovetail with calling convention
  - ▶ Dataflow between gadgets:
    - ▶ Each gadget is memory-memory
- ▶ Turing-complete computation!
- ▶ **Conjecture:** Return-oriented programming likely possible on *every* architecture.



## Conclusions

---

- ▶ Code injection is not necessary for arbitrary exploitation
- ▶ Defenses that distinguish “good code” from “bad code” are useless
- ▶ Return-oriented programming likely possible on *every* architecture, not just x86
- ▶ Compilers make sophisticated return-oriented exploits easy to write

# Summary

- The types of attacks that we must defend against are becoming more complex
- Return-oriented programming shows us that *any* attacker-dictated change in program control flow can lead to arbitrary malice
- Stuxnet shows that ad hoc system defenses can be evaded by an adversary
- We must apply principled approaches to defense to make significant strides in defense